

Prune-Quant Compression Report


Variant	Accuracy	Memory(MiB)	Latency(ms)	Sparsity	Precision

Base Model	10.16%	0.63	1.34	0.00%	fp32
Pruned Model	10.94%	0.25	1.15	60.04%	fp32
Quantized Model	10.16%	0.16	2.77	0.00%	int8
Pruned + Quantized Model	10.94%	0.06	1.78	60.04%	int8


Stage benchmarks (wall time):

evaluate_base	0.16s
sensitivity_scan	0.11s
prune	0.01s
evaluate_pruned	0.08s
quantize_base	0.01s
evaluate_quantized	0.19s
quantize_pruned	0.01s
evaluate_pruned_quantized	0.19s
total	0.75s


Model Comparison: Accuracy, Memory, Latency



Layer Sensitivity Curves





Benchmark — demo-cnn




weight

Base Model — Weight Distribution


weight



weight





Pruned Model — Weight Distribution



weight

Quantized Model — Weight Distribution

weight



weight Pruned + Quantized Model — Weight Distribution weight

