


Prune-Quant Compression Report

Variant	Perplexity	Memory(MiB)	Latency(ms)	Sparsity	Precision
Base Model	89.94	312.47	425.38	0.00%	fp32
Pruned Model	88.31	296.57	537.57	5.09%	fp32
Quantized Model	103.54	78.12	574.18	0.00%	int8
Pruned + Quantized Model	103.61	68.41	573.07	12.42%	int8


Stage benchmarks (wall time):

evaluate_base	29.85s
sensitivity_scan	454.10s
prune	0.91s
evaluate_pruned	39.98s
quantize_base	0.75s
evaluate_quantized	41.46s
quantize_pruned	0.69s
evaluate_pruned_quantized	40.48s
total	608.23s


Model Comparison: Accuracy, Memory, Latency









Layer Sensitivity Curves









Benchmark — distilgpt2-wikitext



Base Model — Weight Distribution




Pruned Model — Weight Distribution




Quantized Model — Weight Distribution


weight




weight




weight




weight



weight



weight



Pruned + Quantized Model — Weight Distribution

